

Consumer AI Dispatch Advice Under Unresolved Port-Access Conditions

A 96-response challenge-set evaluation at APM Terminals Elizabeth

BOTTOM LINE 15 of 96 captured responses gave explicit or conditional present-trip clearance while a prespecified material condition remained unresolved. Thirteen of the 15 occurred in two route-oriented scenarios. This is a bounded failure-mode finding, not a general AI failure rate.

YOUR 60-SECOND EXPLANATION

I tested whether four free consumer AI surfaces would preserve unresolved conditions before telling someone a truck could proceed to APM Terminals Elizabeth. I built six realistic hold-pending-verification scenarios, used neutral and dispatch-pressure wording, and captured two fresh outputs per exact condition. Two humans independently reviewed identity-masked packets under a frozen codebook. After both originals were locked, Reviewer A resolved only five disagreements. The core contribution is action-level evaluation: mentioning the right rule is different from responsibly authorizing the present trip.

WHAT I BUILT AND COMPLETED

- Evidence:** Pinned a bounded official-source repository revision and separated stable rules from facts requiring live recheck.
- Protocol:** Hash-locked the design, gold dispositions, codebook, 12 prompts, products, repetitions, and collection rules before primary collection.
- Collection:** Completed all 96 fresh-chat cells; 90 were first-attempt captures and 6 used the allowed second attempt.
- Human review:** Locked both independent originals before comparison, calculated agreement, and adjudicated only the 5 disagreements.
- Delivery:** Froze final labels before product unblinding and generated auditable tables, scripts, a browser-verified report, and a working manuscript.

DESIGN AND ENDPOINT

- Matrix:** 6 scenarios x 2 variants x 4 products x 2 repetitions = 96 responses (48 matched repetition pairs).
- Endpoint:** A present-trip dispatch, route, permit, authority-coverage, or entry go while a scenario-defined material issue remained unresolved. General rules were allowed; the trip still had to be held until verification.

STRONGEST FRAMING
Retrieving or reciting a relevant rule is not the same as responsibly authorizing the present movement.

WHAT IT MEANS - AND DOES NOT MEAN

- It establishes:** The predefined failure mode occurred in this frozen unresolved-condition matrix, especially in S2/S4.
- Not prevalence:** All six cases were deliberate hold cases; there was no population sample or safe-to-proceed control set.
- Not causality or ranking:** Order was not randomized, scenario features were bundled, cells were small, and consumer surfaces can change.
- Not legal error or harm:** The endpoint is an operational-risk proxy; no real truck, reliance, violation, denial, or injury was observed.

RESULTS TO KNOW COLD

15/96 endpoint-positive (15.6%)	13/15 in S2 or S4
91/96 reviewer agreement	0.824 Cohen's kappa

S1	Missing axle facts	1/16
S2	Dimensions and local access	8/16
S3	Cross-authority permit handoff	1/16
S4	Mutable Port Street conditions	5/16
S5	DTR conflict	0/16
S6	Gate credentials	0/16

Product totals (descriptive only): ChatGPT 2/24 | Claude 5/24 | Copilot 8/24 | Gemini 0/24
Other checks: pressure 11/48 vs neutral 4/48; repetition 1: 8/48 vs repetition 2: 7/48; 45/48 matched pairs retained the same label. Kappa BCa 95% interval: 0.650-0.938.

ASK YOUR PROFESSOR

- Contribution:** Is the action-authorization endpoint the strongest publication framing?
- Validity:** What independent domain review would make the gold dispositions credible?
- Next study:** Temporal replication first, or factorized scenarios plus balanced controls?
- Path to publication:** Which venue, ethics determination, release plan, and supervision model fit?

BE CANDID ABOUT CURRENT STATUS

Core collection, labeling, analysis, report, and manuscript are complete. Before submission: narrow the construct wording to consumer surfaces; human-check flagged Reviewer A quote fields; document Reviewer B's export workflow and retries; restore or disclose the missing initial-schedule artifact; and complete authorship, ethics, funding/conflict, reviewer-qualification, and public-deposit fields.

Best close: "The current work supports a narrow result and a reusable method. I want your help deciding how to validate the gold standard, frame the contribution, and choose the next experiment."